


AI-Assisted Scientific-Support Assessment of Stylized Facts

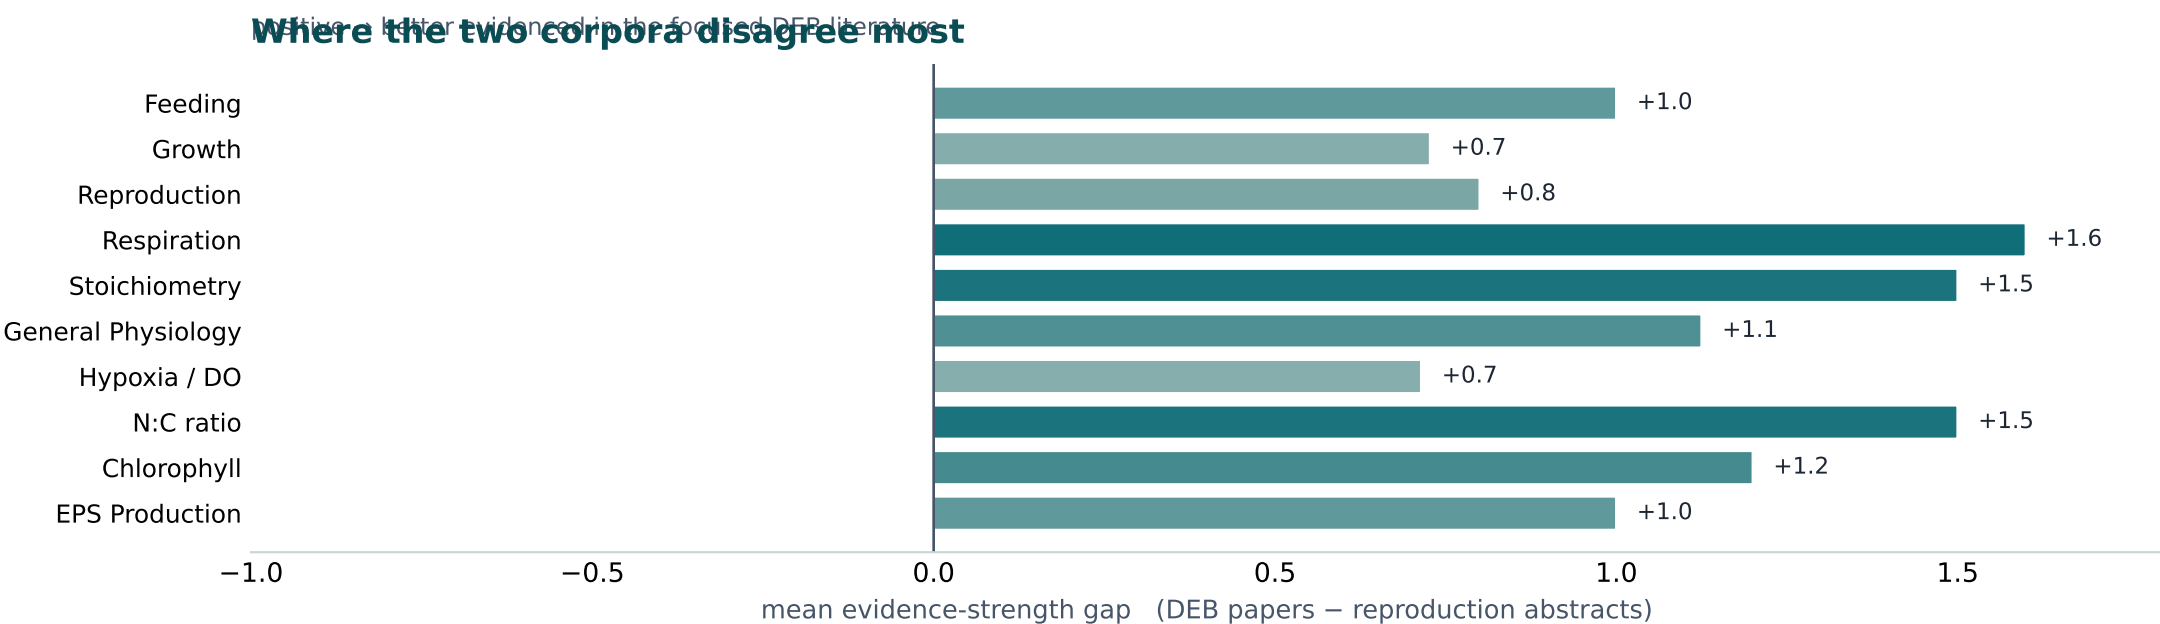
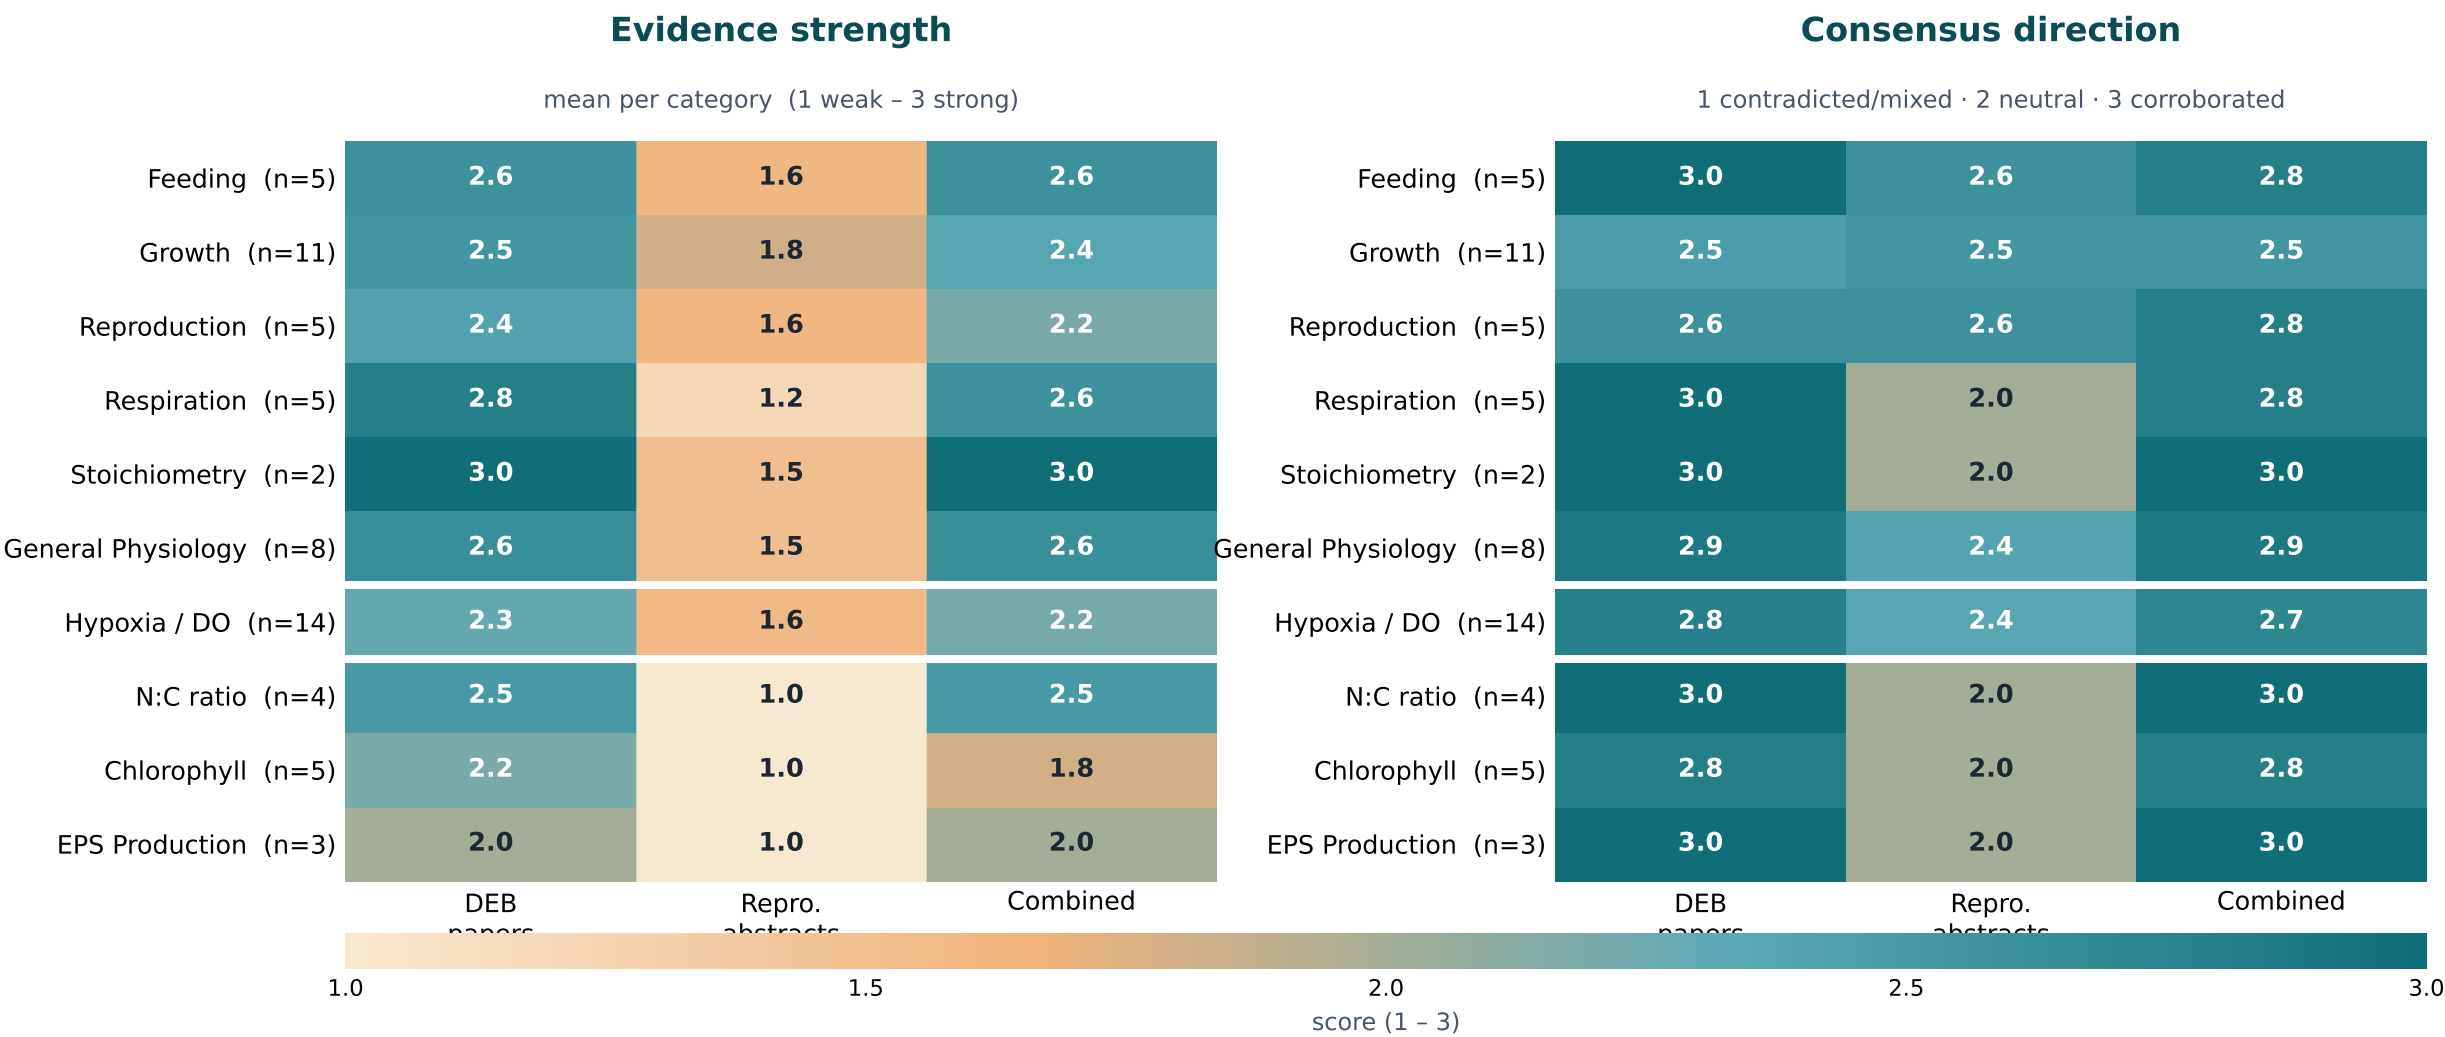
Retrieval-augmented, LLM-adjudicated evidence scoring across two literature corpora

VISUAL
ABSTRACT

How each stylized fact is scored



62	1,643	599,540	2.81 M	3 × 2	372
stylized facts	DEB-domain papers	reproduction abstracts	indexed text passages	scopes × score types	provenance documents



Method in one line

Each stylized fact is embedded and matched against the 60 nearest passages in each corpus; an LLM labels every passage as supporting / contradicting / neutral, and the verdicts are aggregated into two 1–3 scores plus a count of distinct primary references. Every score resolves to its source documents and verbatim snippets.

Caveats: semantic retrieval ≠ proof of support (mitigated by explicit adjudication); reference counts are bounded by retrieval depth, not an exhaustive census; LLM judgments warrant expert spot-checking.