


AI-Assisted Assessment of Stylized Facts — Method & Validation


Full-corpus retrieval-augmented scoring, cross-checked by two independent LLM judges

METHOD



62	1,306	152k	~3.9 M	4.44 M	2
stylized facts	DEB papers	DEB-lit chunks	unique abstracts	abstract chunks	LLM judges

VALIDATION



Full corpus	Cross-checked	Adjudicated	Traceable
~3.9M unique abstracts (99.99% embedded)	two judges, 98–100% within $\pm$ 1	68% of nearest passages filtered as neutral	every score → source passages

Validation philosophy. Similarity locates candidate text; an explicit support/contradict adjudication — auditable to the verbatim passage — decides evidence. Two independent judges (gpt-oss:120b and Claude Opus 4.8) score every fact and agree within one point on 98–100% of facts, so disagreement is a focused signal rather than noise. Because the system reads entire corpora statistically while experts judge from experience, AI-expert divergence is treated as a candidate, objective knowledge gap to investigate — not as automatic AI error.